

VOLLEY Gen5: A System-Level Design Study of a Programmable Electromagnetic CubeSat Deployer

Adityavardhan Mishra

Department of Mechanical Engineering, Symbiosis Institute of Technology
Symbiosis International (Deemed University), Pune, India

Abstract—Rideshare CubeSats commonly inherit the host orbit, while conventional spring deployers provide only a small release impulse. This paper reports Gen5, a fixed system-level computational design for a magazine-fed electromagnetic deployer intended to command the release velocity of unmodified 3U satellites. An ironless double-sided Halbach linear synchronous motor drives a reusable 9.445 kg sled to a 1.5 m release station on 1.8 m longerons; two cassettes hold twelve spacecraft and an eddy-current brake arrests the sled after release. Coupled electromagnetic, circuit, dynamics and mass models predict a 16.029 m/s exit velocity at 10.07 g, 2.78 kJ gross energy draw per shot and 126.6 kg dry mass. Closed-loop simulation predicts 0.0274 m/s (3σ) exit-velocity dispersion. For a stated 450 km, mean-solar-activity case, the orbital model predicts a 28.8 km semi-major-axis increase and a 1.60 lifetime multiplier after one prograde shot; the current rated lifetime result has not been independently rerun. A separate Cartesian propagator confirms the immediate 28.8 km two-body axis change, while a FreeCAD reference assembly exposes an 11 mm side-fed width shortfall and exact track-cassette interference. Independent checks cover selected model quantities, but full-depth thrust, payload compatibility, release-contact dynamics and host integration remain unverified. The design fails its own mass-per-satellite criterion: 10.547 kg versus approximately 2 kg. No article has been built, fired, measured, qualified or flown. Gen6 scaling toward a 1 km/s-class release is future research with no selected architecture.

Index Terms—CubeSat deployment, electromagnetic launch, Halbach array, linear synchronous motor, rideshare, constellation phasing, small launch vehicles.

I. INTRODUCTION

Rideshare secondary payloads may be released into an orbit selected by the primary mission rather than their own target [3]. Conventional spring dispensers supply a small, payload-dependent separation impulse; flight-proven examples include J-SSOD and commercial CubeSat deployers [2], [20]. An operator needing different delivered states can also use host manoeuvres, onboard propulsion, differential drag or an orbital-transport service. A useful deployer comparison must hold payloads, starting orbit, final-state tolerances, deadline, integration burden and failure consequences constant. Repeatable small-satellite fleets are a plausible mission context, but no customer, payload or host has endorsed the Gen5 requirements in this paper.

Electromagnetic release is established prior art, not a novelty claim in itself. Feng *et al.* model a high-acceleration induction launcher for a 20 kg CubeSat [27]; other work studies stacked-spacecraft transfer and reports a transport-mechanism prototype [30], [28], [29]. The question investi-

gated here is narrower: can a shared host-mounted device impart a modelled, individually commanded impulse to an unmodified 3U payload while meeting a chosen acceleration limit and system-level mass, power, containment and mission criteria? The 25 g ceiling is a design choice, not a universal CubeSat qualification limit; the CubeSat Design Specification defers acceptance to the mission provider [23].

This paper evaluates one fixed Gen5 computational configuration. It contributes a modeled machine and shot, selected field and circuit cross-checks, a mass and orbit trade, a reference assembly with an explicit packaging failure, and a register of conditions that prevent a product or flight claim. Host examples are reference calculations, not approved interfaces. The current results include negative design findings rather than a demonstrated customer benefit.

II. RELATED WORK

A. Deployment heritage

The P-POD established the rail-and-spring standard, and the ISS-based J-SSOD and NanoRacks systems industrialized it [2], [3]. The failure record is instructive: the NanoRacks ball-lock anomaly, in which jack-screw preload above 0.11 N·m drove the release mechanism toward seizure, traced to a load path in which ascent preload passed through the release device itself [1]. Zhao *et al.* built a multi-CubeSat electromagnetic transport-and-ejection prototype and report selectable separation speeds of 0.4–0.85 m/s for 3U and 0.67–1.42 m/s for 1U payloads [30]. Programmable electromagnetic release and stacked feed are therefore prior art. Gen5 instead predicts a 16.029 m/s 3U rated shot within a different modeled installed-system design; it has no measured shot and fails its own 3U mass and side-fed packaging criteria.

B. Electromagnetic launch and supporting technology

Coilgun research examined nanosatellite launch at km/s scales [4], [5], establishing contactless acceleration and its costs: the single-stage reluctance machines of that lineage convert 1–2% of stored energy into kinetic energy, though multi-stage induction designs do far better, and the on-orbit CubeSat launcher of [27] reports 14.9–19.9%. Ironless permanent-magnet linear machines occupy the opposite corner of the design space: the Halbach array concentrates flux on one working face [6], and the Inductrack program demonstrated large forces from such arrays with no iron and no sliding contact [8]. Finite Halbach arrays can introduce end-effect thrust

ripple [7], motivating a position-dependent 3-D force check for Gen5 rather than an imported correction factor. Space-oriented supercapacitor studies and component options are reported in ESA-led programmes [9], [10], but no installed Gen5 cell string is qualified. Eddy-current damping has aerospace precedent [11]; the proposed VOLLEY brake has not been tested. On the astrodynamics side, Planet Labs demonstrated propulsion-free constellation phasing by differential drag over campaigns of weeks to months [12], [13]; that timescale is the benchmark against which impulsive seeding is measured. ISRO's POEM programme converted the spent PSLV fourth stage into a stabilized hosted-payload platform across four missions, closing POEM-3 with a controlled, debris-free reentry [14], [15], providing a hosted-stage precedent, not regulatory approval for VOLLEY.

III. SYSTEM ARCHITECTURE

A. Requirements

The driving requirements are: eject an unmodified 3U CubeSat, imposing no armature, plating, or electrical interface on the customer satellite; hold payload acceleration below a *chosen* design ceiling of 25 *g*, a requirement this work sets on itself rather than one any standard imposes; mount on a restartable upper stage, kick stage, or hosted orbital platform through a standard ESPA bolt pattern, a host class whose accommodation envelope is not public and against which the closed 1839 mm length therefore cannot yet be shown to fit (Sec. XV); provide per-satellite programmable velocity with dispersion small against its astrodynamics effect; carry a twelve-satellite manifest with autonomous sequencing; provide inhibits and a defined abort; and hand recoil to the host in a form its attitude control absorbs at negligible propellant cost.

B. Derivation of Principal Design Parameters

The values that define the machine combine internal requirements, assumptions and externally sourced limits; Table I records the governing relation for each. Three deserve explicit statement. The acceleration ceiling is chosen, not derived. An earlier version of this work obtained it by integrating the NASA GEVS [21] protoflight random-vibration spectrum to 14.1 *g*_{rms}, taking a 3 σ envelope of roughly 42 *g* as the launch load, and setting a 25 *g* cap beneath it. That derivation is withdrawn. Broadband random vibration and a 162 ms commanded acceleration are different load cases: *g*_{rms} is the root-mean-square of a stationary random process over a spectrum, and multiplying it by three produces a peak-response estimate for that process, not a quasi-static limit load a structure is qualified to carry. Neither GEVS nor the CubeSat Design Specification [23] publishes a universal quasi-static level for CubeSats; the CDS defers test levels to the launch provider, and levels therefore vary by vehicle and by mission. 25 *g* is retained here as an internal design requirement, the number the retention chain, the cradle, the abort logic and the payload-family table are all sized against, and it is not offered as evidence that any satellite is qualified for it. An earlier version of this design was acceleration-limited: at a 4.9 kg parametric

TABLE I
GOVERNING RELATION FOR EACH PRINCIPAL PARAMETER

Parameter	Value	Set by
Accel. ceiling	25 <i>g</i>	<i>chosen</i> design limit (Sec. III-B)
Track length	1.5 m	$v = \sqrt{2aL} + \text{coast zone}$
Rail material	7075-T6, hard anodize	CDS rail spec [23]
Bank voltage	96 V	SiC 1200 V derated 60%
Bank capacitance	6 F	droop eqn (2), $\leq 5\%$ SoC sag
Sled magnets	N45SH	thermal $\times$ field (Sec. V-F)
Brake cap	200 <i>g</i> sled	magnet bond MoS (Sec. III-F)
Abort point	45% stroke	braking-distance equality
Magazine	12 (2 $\times$ 6)	cassette CoM symmetry

sled the ejection ran at 16.3 *g*, close enough to the cap that the cap set the velocity through $v_{\max} = \sqrt{2a_{\max}L}$. That is no longer the binding constraint. With the sled computed at 9.4 kg from the CAD solid volumes the same commanded force accelerates a heavier moving mass, so the ejection sits at 10.1 *g* and the 1.3 m zone yields 16.0 m/s. The machine is now *thrust-and-mass* limited rather than acceleration-limited, with a factor of 2.5 of margin against its own ceiling left unused. Recovering velocity therefore means removing sled mass or raising sheet current, not shortening the stroke; a 0.2 m coast-trim zone completes the 1.5 m axis. Stroke buys velocity only as $\sqrt{L}$, so the trade is unattractive: total track lengths of 1.0, 1.2, 1.5 and 1.8 m give 12.6, 14.1, 16.0 and 17.8 m/s respectively, and the envelope is already the binding packaging constraint (Sec. XV). The capacitor bank follows from the shot energy and the droop it actually reaches: at the 2.78 kJ draw the selected 6 F bank sags to $V_1 = 0.947V_0$, and

$$C = \frac{2E}{V_0^2 - V_1^2} = 6.00 \text{ F} \rightarrow 6 \text{ F}, \quad (1)$$

realized as thirty-two series 190 F, 3 V commercial cells for a 96 V rail compatible with 1200 V SiC devices at a comfortable derating.

C. Topology Trade

The payload settles the trade before the electromagnetics do. Exit velocity is bounded by

$$v_{\max} = \sqrt{2a_{\max}L}, \quad (2)$$

so at the 25 *g* design ceiling over a 1.3–2 m stroke the reachable range is 25–31 m/s. The coilgun's defining advantage, a km/s-class ceiling, is outside this Gen5 payload load envelope, while its costs remain: microsecond timing against the suck-back effect, an armature bolted to the customer's satellite, and no abort once fired. Efficiency is deliberately not among them: at 14.9–19.9% [27] a multi-stage coilgun is comparable to the figure reported here, and an earlier version of this argument claimed otherwise on the strength of a single-stage number. A linear synchronous motor concedes nothing in the reachable envelope and inverts every cost: high drive efficiency, continuous servo control, and a reusable sled that carries the magnets so the satellite carries nothing. Table II summarizes.

TABLE II

TOPOLOGY TRADE. THE VELOCITY ROW IS EVALUATED AT THIS DESIGN'S CHOSEN 25 g CEILING, NOT AT A QUALIFICATION LIMIT OF THE PAYLOAD

Criterion	Reluctance coilgun	Ironless LSM + sled
Velocity at the 25 g ceiling	25–31 m/s	25–31 m/s
Electrical→payload efficiency	14.9–19.9% [27]	18.8% (Sec. V)
Satellite modification	armature required	none
Velocity control	pulse timing	closed-loop servo
Abort	none once fired	to ~45% stroke
Field on stored satellites	pulsed, unshielded	Halbach self-shielded

D. Layout and Firing Cycle

The model places the release station 1.5 m from the breech on 1.8 m structural longerons; the powered stroke is 1.3 m, followed by a 0.2 m coast-and-trim interval (Figs. 1 and 2). The stator is a double-sided ironless three-phase winding. The modeled mover is a 9.445 kg sled carrying opposed Halbach arrays (48 mm wavelength, 8 mm N45SH blocks, four blocks per wavelength) across a 12 mm winding gap. The two six-satellite cassettes are conceptual feed elements. Their current side-fed reference placement fails the width and exact-solid interference checks reported in Sec. XV.

The intended sequence is follower advance, restraint engagement, motor acceleration, sled braking and payload departure. Feed forces, contact transitions, retention during ascent, release timing and tip-off have not been verified for an actual satellite. The current model assumes that the payload remains seated against the aft backstop during acceleration and separates when the sled decelerates; that assumption requires moving-contact analysis and instrumented tests. The ideal internal mass-motion model gives no final rigid-body rate after the sled stops, but it predicts a transient return rate of 0.130 deg/s for the assumed 500 kg host and 0.385 deg/s at 200 kg; both declared peak-rate bands fail. The earlier 8.2 s rate-null and 18.1 s cadence floor are superseded. An assumed 166 mm lever arm gives a 0.40 deg attitude offset in the 500 kg example. Controller restoration time and structural ringing have not been established. The adopted mission-planning interval is 1200 s between shots; 10–20 s is only a modeled recharge floor, so twelve scheduled shots span approximately 4.0 h and 2.6 orbits if all other operations succeed.

E. Restraint, Inhibits, and Abort

Launch restraint separates the preload path from the release path, the inverse of the NanoRacks fault [1]. One-shot retention gates at each cassette exit carry ascent preload directly into structure; the escapement is caged during ascent; the sled is held by a motorized over-center cam lock. Firing requires a three-inhibit chain: redundant seat switches, an attitude-valid flag from the host, and a sequencer arm. Abort is available to approximately 45% of the stroke, beyond which braking distance for the combined mass exceeds remaining track; the muzzle buffer is sized to capture an aborted combined mass. The retention gate carries the ascent preload of a six-satellite




Fig. 1. System block diagram: power, control, mechanism, and host-interface chains.




Fig. 2. Plan-view layout against an ESPA-Grande-class envelope (not to scale), retained as a dimensional reference. The closed envelope exceeds that class in length by 44%; ESPA-Grande port compliance is not a requirement of this design, and the target host class is discussed in Sec. XV.

stack through two D9 A-286 shear pins of 41.0 kN double-shear capacity; the escapement, caged behind the gate, sees none of this. The quasi-static case taken at this design's own 25 g ceiling, $24\text{ kg} \times 25\text{ g} = 5.9\text{ kN}$, is *not* the sizing case, and the gate was originally sized against it with two D6 pins. Miles' equation on the GEVS protoflight spectrum ($14.1\text{ g}_{\text{rms}}$, $0.16\text{ g}^2/\text{Hz}$) at the track's 109 Hz fixed-fixed mode gives a 3σ load of 11.7 kN at $Q = 10$ and 20.2 kN at $Q = 30$, two to three and a half times the quasi-static figure, which left the D6 margin negative at $Q = 30$. The pins were not undersized; the load case was. Resizing to D9 raises the margin of safety to +0.45 at $Q = 30$ and keeps it positive across $Q = 10$ –30 for an added mass of 11 g, so the retention chain no longer depends on where the unmeasured structural damping lands, which matters because the same unmeasured Q governs the force-ripple resonance discussed in Sec. XV and cannot be resized away there. Random vibration, not quasi-static preload, is the driving case for the retention chain, and is now the case it is sized against. The abort threshold is where the sled braking distance under the same thrust equals the remaining track, which for the loaded 13.4 kg carriage falls at 45% of stroke and 11.1 m/s. The commit fraction is unchanged by the heavier sled, braking distance and travelled distance scale with




Fig. 3. Eddy-brake arrest, taper-limited to 200 g on the sled.

mass identically, so the mass cancels, but the velocity at the commit point falls with the acceleration.

F. Arrest, Power, and Materials

The sled leaves release with about 1.21 kJ. Motor regeneration alone cannot arrest it, because braking force is bounded by the same thrust constant as acceleration; that bound sets what regeneration can and cannot do, and both halves matter. Braking at the rated sheet current over the 39 mm of stator the closed envelope allows downstream of release extracts 51.8 J of mechanical work for 0.4 J of copper, returning 47 J to the bank, and the sled still reaches the brake at 15.7 m/s carrying 1162 J. A copper-fin eddy-current brake therefore provides bulk arrest: force proportional to velocity, contactless, with flight heritage in the damper class [11]. Authority is abundant, so the pole entry is tapered to cap sled deceleration near 200 g. This limit is set by the magnet retention: a 65 g Halbach block at 200 g loads its bonded interface (10.8 cm²) to only 0.12 MPa against a conservative 10 MPa epoxy allowable, a margin of 84, so 200 g is a deliberate conservatism protecting the brittle sinter rather than a strength limit. The 0.344 kg copper fin absorbs the sled's residual 1162 J per shot at an adiabatic 9.0 K per-shot transient rise; cooling between shots has not been modelled; a ring-spring catches the residual below 1.5 m/s (Fig. 3). Pulse energy comes from a 6 F, 96 V supercapacitor bank switched by a SiC bridge [9]. Materials follow three rules: nothing conductive rides the changing field (titanium sled chassis, PEEK formers), nothing near the track is ferromagnetic except by design (BeCu springs, A-286 fasteners), and nothing outgasses near customer optics (E595-compliant polymers, dry films, sealed roller cartridges).

IV. MODELS AND VERIFICATION

A. Field Model and Independent Verification

The Halbach airgap field is a decaying wave, $B(y) = B_0 e^{-ky}$ with $k = 2\pi/\lambda$, summed over the two opposed faces. Because this carries the force budget, it was checked against an independent magnetostatic computation (analytic cuboid superposition, magpylib), with the array rotation-sense convention determined empirically by probing a single array on both faces; the exercise caught two successive sign errors whose outputs matched single-array fingerprints. Converged

results: single-array mid-gap 0.351 T against the wave model's 0.351 T; double-sided mid-gap peak 0.694 T against a predicted 0.702 T. The stray field behind the array back face, which sets the magnetic keep-out for stored satellites, computes to 22.7 mT at 10 mm, 4.3 mT at 20 mm, and 0.4 mT at 50 mm.

B. Winding-Resolved Thrust Constant

The airgap field is sampled in all three dimensions before the Lorentz sum: B_y is Gauss-Legendre averaged across the array's full 90 mm depth rather than taken on the centre plane. That distinction is worth stating because the centre-plane value is 4.4% higher, and using it, as this work did until an independent depth-resolved integral was run, overstates the thrust constant by the same margin and the exit velocity by 2.2%. A slotless three-phase belt winding (pole pitch $\lambda/2$, belts of $\lambda/6$, winding factor 0.955) was resolved against the verified field by direct Lorentz integration with position-locked sinusoidal commutation, the field translated with the sled. The result is a thrust constant

$$K_t = 10.54 \text{ N/(kA/m)}, \quad \text{ripple } \pm 1.01\% \text{ (6th harm.)}, \quad (3)$$

shown in Fig. 5. Fig. 4 shows the field the integral is taken over, and the reason the integral must be taken over it: sampled on the array's centre plane B_y peaks at 0.5041 T, while the mean across the 90 mm depth is 0.4759 T, a ratio of 0.9440. Treating the centre-plane value as though it held across the depth is the 4.42% overstatement Sec. XV records.

Synchronous extraction carries the one-half traveling-wave factor that lumped $\langle B \rangle KA$ models hide; the corrected constant is recovered operationally by rating the sheet current at 140 kA/m and commanding 126 kA/m (21 A/mm² pulse-only in the copper), where adiabatic heating remains 0.35 K per 162 ms shot: the shot's 854 J of copper loss into the 6.29 kg of stator copper at 385 J/(kg·K), quoted to two significant figures throughout as 0.3527 K computed. An earlier version of this section rounded it to 0.3 K while Sec. IX carried 0.35 K; the two are the same number and the paper now states it once.

C. Shot, Servo, and Astrodynamics Models

The shot integrates sled-plus-satellite dynamics against the supercapacitor circuit with explicit copper loss and a 95% converter. Velocity control uses a position-scheduled profile with proportional correction and a coast-trim residual correction from photogate measurement; an 800-run Monte Carlo over thrust-constant, mass, force-ripple, and sensor tolerances gives the dispersion. The proportional gain is not tuned to that dispersion but designed against stability margins, as Sec. IV-D sets out. Orbital lifetime uses per-revolution orbit-averaged decay (Gauss tangential equations, quadrature over eccentric anomaly, static exponential atmosphere at mean activity [17]); the method was cross-validated against an independent Cowell RK4 propagation with drag, the two agreeing on 30-day semi-major-axis decay to 99.4%. Constellation seeding uses $\Delta a = 2a \Delta v/v$ and the resulting period difference; deployment safety propagates the stage and all deployed ellipses with




Fig. 4. Halbach airgap field. Left, B_y in the x - y plane over one wavelength; right, the profile through the array's 90 mm depth. Drawn by the same field builder the thrust integral uses, so the two cannot disagree.




Fig. 5. Winding-resolved thrust over one wavelength of sled travel.

Kepler-plus-secular- J_2 motion for 30 days, with candidate minima refined at 0.25 s sampling.

D. Velocity-Loop Stability

The control law divides the commanded sheet current by the modelled thrust constant and multiplies by the modelled mass, so the plant's own K_t/m cancels and the open-loop transfer reduces to $L(s) = K_p e^{-s\tau}/s$. The proportional gain is therefore not a current gain but an acceleration per unit velocity error, in s^{-1} , and *its numeric value is the gain crossover frequency in rad/s*. The total lag τ is the sensing and computation transport delay plus the half-sample of a zero-order hold; both are pure delays and add. No sensor has been selected, so the delay is a stated assumption, swept rather than asserted, and 0.6 ms is used as the reference value with a 0.1 ms hold at 5 kHz.

That reduction sets the design constraint. An earlier gain of $3500 s^{-1}$ places the crossover at 557 Hz, five times the track's 109 Hz first bending mode and above the 48 Hz mode as well, giving -50.4° of phase margin and -3.9 dB of gain

margin at the reference delay, with marginal stability reached at only 0.45 ms of total lag (Figs. 6 and 7). Two properties of the simulation had concealed this: the loop fed back an undelayed state, which is the one condition under which that gain is stable, and the command is clipped at the winding rating, so the linearly unstable loop behaves as a bang-bang relay whose mean follows the feedforward term while the terminal photogate trim removes the residual. The reported dispersion was consequently governed by the saturation limits and the terminal correction rather than by the feedback it was attributed to, which is why the closed-loop figure is reported here only alongside the margins that make it a property of a design.

The gain is instead sized as the largest value simultaneously holding $\geq 50^\circ$ of phase margin at the reference delay and keeping closed-loop bandwidth at or below one third of the first bending mode, which gives $195 s^{-1}$: crossover 31.1 Hz, phase margin $+82.2^\circ$, gain margin $+21.2$ dB, and no excursion above the current rating anywhere in the stroke. The gain falls by a factor of eighteen and the 3σ dispersion does not move, 0.0271 to 0.0267 m/s. The loop never required the bandwidth: dispersion is set by the terminal trim and by the thrust-constant and mass tolerances, not by loop gain. The binding constraint is now the structural one, so a measured first mode below 109 Hz would lower the designed gain with it. Classical margins on a 162 ms trajectory are necessary rather than sufficient, the plant is treated as rigid, and the structural modes enter only as a frequency the bandwidth is held away from; a compliant-track model remains outstanding.

V. PERFORMANCE RESULTS

A. Launch Performance

At the rated point (3U reference, 9.445 kg sled input from historical Gen3 CAD solid volumes, 1.3 m acceleration): exit velocity 16.0 m/s at 10.1 g, pulse 162 ms, peak current 320 A,




Fig. 6. Open-loop velocity-loop response at the earlier and the designed gain, with the 48–109 Hz band of track modes shaded and both gain crossovers marked. Phase depends only on frequency and delay, so the two gains share one phase curve.




Fig. 7. Phase margin against transport delay. The earlier gain crosses into instability at 0.35 ms of transport delay; the designed gain holds more than 55° across the whole swept range.

and bank state-of-charge sag 5.2% (Fig. 8). Energy drawn is 2.78 kJ per shot, with copper heat of 854 J and a further 76 J dissipated in the bank's own series resistance. Regenerative braking returns 47 J after release, so the net draw is 2.74 kJ. The payload receives 514 J, a 18.8% electrical-to-payload efficiency; the sled's share is split between 47 J recovered and 1162 J dissipated in the arrest brake. Efficiency falls with the heavier sled for two compounding reasons: a larger fraction of the same mechanical work goes into a mass that is then braked away, and the longer 162 ms pulse accrues more copper loss at unchanged current density. Recovery is what partially offsets the first of those, and it is bounded by geometry rather than by loss: over the available 39 mm the braking pulse lasts 2.5 ms and burns 0.4 J of copper, against 854 J over the 162 ms shot, so recovered energy rises monotonically with braking force up to the winding's rating and with distance beyond it. A 500 mm




Fig. 8. Shot profile at the rated point: (a) velocity with coast-trim zone; (b) bank voltage and current.




Fig. 9. Closed-loop exit-velocity dispersion, 800 Monte Carlo runs.

section would return 540 J and a 1 m section would arrest the sled outright while returning 727 J, but the closed envelope already exceeds its target class by 44% (Sec. XV), so distance is not available here. Peak current during recovery is 212 A, below the shot's own 320 A, so the drive is not re-rated. The closed-loop Monte Carlo returns a dispersion of 0.0274 m/s (3σ) at a 15.8 m/s fleet setpoint (Fig. 9), dominated by the sensing chain rather than plant tolerances. The setpoint must sit below the open-loop ceiling for the servo to retain authority; at 98.85% of it, the nominal headroom is 0.188 m/s. Table III gives the payload family: the system is force-limited above 1U, so heavier payloads trade velocity, never structural margin, and the deployer mass each customer pays for varies by a factor of 109 across the family (Sec. V-B).

A separate algebraic audit in the companion repository recalculates the payload kinetic energy, constant-acceleration equivalent stroke, dry-mass sum and brake-work identities from the captured JSON. These agree within the declared rounding bands. Subtracting the reported payload and sled kinetic energies, shot copper heat and bank ESR from gross electrical draw leaves 124.488 J (4.47%) unitemized. This is an accounting gap, not an identified physical loss; a complete electrical budget and supplier-backed installed bank remain open.

B. Payload Class and the Mass-per-Satellite Argument

Table III is read most usefully from its last column rather than its first. Exit velocity is nearly flat across the family, because the 9.445 kg sled is most of the moving mass: re-

TABLE III
PAYLOAD FAMILY AT THE RATED POINT (1.3 M ACCELERATION, 9.445 KG CAD-DERIVED SLED). DEPLOYER MASS PER SATELLITE USES THE 126.6 KG DRY MASS AND A PACKING MODEL CALIBRATED SO THE 3U CASE RETURNS THE TWELVE THE MAGAZINE IS LAID OUT FOR. EVERY COLUMN IS REGENERATED FROM `PAYLOAD_FAMILY.JSON` AT THE CURRENT OPERATING POINT; THE VALUES PUBLISHED UNTIL 2026-08-20 PREDATED BOTH ADR-030 AND THE ENCLOSURE BUILDUP

Class	Exit vel.	Accel.	Per load	kg/sat.
PocketQube 1P (0.25 kg)	18.9 m/s	14.0 g	326	0.388
ThinSat (0.28 kg)	18.8 m/s	13.9 g	123	1.029
PocketQube 3P (0.75 kg)	18.4 m/s	13.3 g	108	1.172
TubeSat (0.75 kg)	18.4 m/s	13.3 g	41	3.088
1U (1.33 kg)	17.9 m/s	12.6 g	40 [†]	3.165 [†]
3U (4 kg)	16.0 m/s	10.1 g	12	10.550
6U (8 kg)	14.1 m/s	7.8 g	6	21.100
12U (12 kg)	12.7 m/s	6.3 g	3	42.200

moving the payload entirely buys 19% of velocity, so velocity is a property of the sled and the stroke and not of the customer. Shortening the magnet array does not recover it either, since array length sets thrust and sled mass together, a 150 mm array gives a 5.15 kg sled but a 4.65 N per kA/m thrust constant, and the 3U case falls to 12.9 m/s.

The column that does move is deployer mass per satellite: by a factor of 27 between the 3U case and PocketQube 1P, and by 109 across the full column. At 3U the machine spends 10.55 kg of deployer per customer, which loses to a cold-gas propulsion module delivering the same 16.0 m/s at 0.5–1.2 kg by a factor of between nine and twenty-one, 12.4 against the 0.85 kg mid-class figure. At PocketQube 1P it spends 0.388 kg and wins by 1.3 to 3.1. The comparison inverts on payload class, and on no machine parameter at all: not the thrust constant, not the stroke, not the sheet current. It inverts less decisively than it did: the figures published until 2026-08-20 were 7.04 and 0.24 kg, and the enclosure buildup that replaced an 8.00 kg packaging placeholder with 50.04 kg of derived line items moved both.

Three qualifications belong with those counts rather than beneath them. First, they are volumetric. The packing efficiency is calibrated against the one configuration that has been laid out in CAD, the 3U case, where a raw volume ratio predicts 21 and the magazine holds 12, and the resulting 56% is then applied to classes for which no cassette exists. Second, the feed engages the CubeSat Design Specification corner-rail interface at the cradle, the escapement and the retention gate; PocketQubes, ThinSats and TubeSats do not have that interface, so a smaller class needs its own cassette, cradle and gate, which is mechanical design rather than a parameter change. Third, hundreds of satellites per load is a different machine from twelve: the campaign thermal case, the bank recharge duty and the escapement cycle life were all sized for twelve shots, and the bank is already the binding constraint at that number (Sec. XV). What Table III establishes is that payload class dominates the economic argument, not that a small-payload variant has been designed.

The comparison the 3U column loses most clearly is against

the incumbent it was meant to match. A canisterised spring dispenser occupies roughly 6 kg per 3U slot, and an acceptance band written before the comparison script existed asked for parity within 15%. The measured ratio is 1.758: this deployer is 76% heavier per satellite than the hardware it would displace, and the band fails. It failed on a change to this work rather than to the comparator, a packaging rollup that had carried an 8.00 kg placeholder was replaced by 50.04 kg of derived line items, and parity did not survive it. The project’s own kill criterion on mass per satellite, fixed at an estimated 2 kg, is crossed by $5.3\times$. No enclosure change reaches parity: the lightest alternative structure costed, a honeycomb sandwich, would give 8.87 kg per satellite and a ratio of 1.48. The route that moves this criterion is a smaller payload class, which is what the left of Table III prices and what Sec. XV records has not been designed.

It is not for want of looking. Every kilogram of the dry mass was attributed to the requirement that causes it, six requirements, from *the satellite is unmodified* to *twelve satellites share one drive*, and each of the 64 subsets was then deleted in turn to find the lightest machine any combination of relaxations could reach (Fig. 10). Three of the six carry no mass on their own. The best corner deletes both of the two that do, giving up the reusable mover and the requirement that the energy arrive during the shot, and it removes 37.89 kg, or 29.90%. 88.67 kg survives every deletion in all 64 corners, 70.06% of dry mass, 7.39 kg per satellite against a 2 kg criterion. Deleting requirements does not close it; only carrying more, smaller satellites does. The ledger is an attribution and an upper bound on what relaxation buys, not a sizing model, and it does not claim the surviving mass is irreducible by better design.

The architecture does not meet the spring-dispenser mass parity criterion. A separate earlier claim that waiting between releases gives satellites a persistent phase offset at zero relative velocity is also withdrawn: co-orbital zero-impulse releases from an unchanged host share the same state at a common epoch, as Sec. V-C explains. The contribution evaluated here is a modelled, commanded per-satellite orbital impulse without modifying the satellite, at a design acceleration of 10.1 g. Payload structural compatibility remains unverified; its qualified load environment and a host integration review are required. The modelled deployer costs 1.758 times the mass per satellite of the spring comparator.

C. Astrodynamical Utility

A prograde ejection raises apogee while leaving perigee at the deployment altitude, so the correct claim is a longer-lived ellipse. At 450 km with ballistic coefficient 61 kg/m², one 16.0 m/s shot multiplies lifetime by 1.60 at mean solar activity (base lifetimes 0.52–2.61 yr across the activity range; Figs. 11 and 12).

An earlier version of this work claimed the multiplier was invariant to two decimal places across ballistic coefficients of 40–90 kg/m² and across a fivefold density range, and offered that invariance as the result the analysis could defend.




Fig. 10. Constraint ledger: dry mass attributed to the requirement causing it, and the floor reached by deleting requirements in all 64 combinations. The floor is 88.67 kg, 70.06% of the 126.56 kg rollout.

That claim does not survive an independent propagator. Run headless against MSISE90 with 20×20 gravity, lunisolar third bodies and solar radiation pressure, GMAT R2022a reproduces the multiplier at mean and high activity but returns 2.07 at low activity, a spread of 18% against a 5% acceptance band fixed before the run. The reason is a defect in how the sweep was posed rather than in the arithmetic: the model represents solar activity as a uniform multiplicative scale on density, and ballistic coefficient enters the drag term in the same multiplicative slot, so both sweeps divide the two lifetimes being compared by the same factor. Their ratio cannot move, whatever the atmosphere does. A real atmosphere changes the *shape* of the density–altitude profile with activity, the boosted orbit’s apogee samples that profile some 37 km higher than the baseline’s, and the ratio then does move. The multiplier is therefore quoted here at a stated activity level, and no invariance is claimed.

Differential ejection velocities put adjacent satellites on different periods: splits of 2, 5, and 10 m/s reach 30° spacing in 6.9, 2.8, and 1.4 days. Under identical density and geometry assumptions, differential drag between 3:1 area configurations achieves the same spacing in 25.0 days (Fig. 13), consistent with flown drag-phasing campaigns operating over weeks to months [12].

A common-epoch comparison must account for the actual state imparted at release. In a two-body reference case, a satellite released with zero relative velocity from an unchanged host follows the host orbit. A second identical release later follows that same trajectory; at any common time the two satellites are co-located. The host’s angular travel between release events is not persistent separation between the satellites. The 1200 s inter-shot cadence is therefore a thermal and recharge assumption, not a phasing result. Conventional spring impulses, host manoeuvres and subsequent differential drift remain valid comparators and must be included in any mission-level phasing trade.

A differential deployment impulse can change orbital period

and create a growing phase separation, as the preceding model shows. It does not establish a permanent fixed angular offset without later control or drift management. The differential-drag comparison in Fig. 13 is conditional on the stated area and atmosphere assumptions. A matched comparison of commanded electromagnetic shots, spring deployment, host manoeuvres and drag is still needed before asserting a phasing advantage.

A current-rated cross-check independently propagates the Cartesian post-release state for two orbital periods and recovers a 28.800775 km semi-major-axis rise at 16.029 m/s. The 0.02 m solver/identity band passes. This verifies instantaneous two-body geometry, not the 1.60 lifetime multiplier, host campaign or launcher performance. The method and exact inputs are archived in the companion repository.

The Gen5 model instead evaluates a commanded change in the deployed satellite’s orbit. Its 16.0 m/s prograde shot predicts a 28.8 km semi-major-axis increase and a 1.60 lifetime multiplier at the stated mean-activity condition. A 2.5 m/s spring is also an orbital impulse and produces an 8.2% modelled lifetime extension under the same setup. Release timing without an impulse changes neither orbit; a manoeuvring host can change the initial orbit of later releases. These are different mission cases, and the quoted Gen5 benefit is model output rather than demonstrated performance.

D. Deployment Safety

Every prograde shot leaves an ellipse whose perigee sits at the firing altitude, so the fleet re-crosses the stage’s orbit by construction, and the along-track phase between a deployed satellite and the stage realigns on a period of 10.3 days at the rated velocity (Fig. 14). The closest approach reached within a 30-day screening window is deliberately *not* quoted here as a safety figure: because both bodies re-cross the same altitude, that minimum is fixed by the exact along-track phase at crossing, which is acutely sensitive to ejection point velocity, in the model a $\pm 2.5\%$ change about the rated point sweeps




Fig. 11. Orbital lifetime vs deployment altitude (BC = 61 kg/m², mean activity).




Fig. 12. Solar-activity sweep against GMAT. The static-atmosphere model returns a flat ratio by construction, not by physics; an independent propagator does not.

the 30-day minimum by more than an order of magnitude. It is a near-resonant sample of the beat geometry, not a design property. Two statements survive that sensitivity and carry the safety argument. First, per-shot collision-avoidance (COLA) screening is mandatory for every ejection, because the approach geometry cannot be certified from the nominal velocity alone. Second, disposing of the host stage before the first 10.3-day realignment, routine for a restartable stage, removes the recurring close-approach mechanism at its source, before the fleet can return to the stage's along-track position. Per-shot conjunction products remain a mission-operations deliverable.

E. Error, Tip-Off, and Mass

The 0.0274 m/s dispersion maps to ± 0.10 km apogee placement at 450 km: sub-kilometre deterministic insertion for a satellite with no propulsion. Tip-off is modelled rather than bounded, and the risk is not where the analytic budget put it. Release occurs 200 mm, 12.2 ms, into the coast-and-trim zone with commanded force already at zero, so at the residual latch and friction load the mechanism tolerates roughly 250 μ s of asymmetry and still lands two orders of magnitude inside the 2 deg/s comparator. What is not benign is the start of the stroke: the payload centre of mass sits 70 mm off the thrust line, the 413 N push therefore applies a 28.9 N·m moment, and the payload accelerates across its cradle clearance at




Fig. 13. Modelled velocity-split phasing versus differential drag under identical assumptions. The figure does not include spring impulses, host manoeuvres or a complete mission-level comparator.




Fig. 14. Satellite-stage range over 30 days, staggered firing at the rated velocity.

688 rad/s², arriving at 36–231 deg/s depending on fit, 18 to 115 times the comparator. Whether that impact has damped out by release depends on a restitution model this study does not have, so tip-off is stated as a requirement on a cradle preload exceeding the 85 N couple reaction and on a release residual below about 1 N, against a release mechanism that remains undefined (Fig. 15). Recoil per shot is the satellite's momentum only, 64.1 N·s, the sled's share returning through the brake. The solid model closes dry mass at 126.6 kg (loaded 174.6 kg, centre of gravity 0.46 m from the breech), with the ironless stator at 7.3 kg of copper and formers; the system uses roughly a third of an ESPA Grande mass allocation. It does not currently fit the corresponding envelope: the closed CAD assembly spans 1839 mm along the track against the roughly 1270 mm of that class, because the arrest brake sits beyond the release point and the enclosure must span it. The third of those routes is the one taken: the target host class is a restartable upper stage, kick stage, or hosted orbital platform, and ESPA-Grande port compliance is not a requirement of this design. That is a re-scope, not a fix, and its cost is stated rather than absorbed. Shortening the track to 1270 mm would leave a 731 mm acceleration zone and, since velocity goes as $\sqrt{L}$, would cut ejection from 16.0 to 12.0 m/s, a 25% loss in the quantity every performance claim here rests on. The re-scope avoids that and gives up the ESPA-Grande port population in




Fig. 15. Tip-off error budget against deployer requirement classes.




Fig. 16. Payload family: force-limited above 1U; velocity, not safety, is traded. [†]Counts are *volumetric*, calibrated so the 3U case returns the twelve the machine is laid out for. A subsequent fixed-cell design study replaces them where a cell geometry actually exists: 1U gives 36 per load at 2.125 kg, not 40 at 1.92, because three 100 mm units plus two dividers leave 37.5 mm of a 340.5 mm cell unusable, and ThinSat and 12U are refused outright because two of their dimensions exceed the 100 mm cell section. The binding constraint is the 166 mm cassette width.

exchange. It also leaves the packaging question unanswerable for now rather than answered: no accommodation envelope for the target host class is public, so the 1839 mm length cannot presently be shown to fit anything, and obtaining that envelope is part of the same single data exchange Sec. VI identifies. This is an open item, and the honest description of it is that a measured failure against a published limit has become an unmeasured unknown against an undisclosed one.

VI. GENERALIZED HOST INTEGRATION

A. Interface Requirements for Any Host

The reference host analysis assumes a mechanical mount, 150–300 W recharge allocation or an equivalent battery, an attitude reference and control authority for shot transients, and a compliant disposal plan. These are requirements to negotiate, not verified provisions of any particular stage. The quoted full-manifest momentum and propellant estimates are idealized budget calculations; actual correction, attitude recovery and disposal cost depend on host dynamics, thruster performance and the delivery sequence. Table IV varies the assumed host mass to show the translation sensitivity. A free-flyer variant would need its own GNC and is outside this Gen5 study.

TABLE IV
PER-SHOT RECOIL (64.1 N·S) VERSUS HOST CLASS

Host class	ΔV per shot	12-shot total
200 kg free-flyer	328 mm/s	3.93 m/s
300 kg kick stage	219 mm/s	2.62 m/s
600 kg kick stage	109 mm/s	1.31 m/s
900 kg (PS4 class)	72.8 mm/s	0.87 m/s
420 t (station class)	0.2 mm/s	2.3 mm/s

B. Worked Host Examples

ISRO’s POEM missions demonstrate that a spent stage can support hosted experiments and controlled disposal [14], [15]. Published Skyroot descriptions indicate interest in restartable orbital operations [18], [19]. These are contextual examples only. No host has supplied an interface control document, confirmed available mass/power/attitude margins, accepted a magnetic payload environment or approved VOLLEY accommodation. A real host trade must obtain the stage mass and inertia, permitted envelope, bolt pattern, power/command provisions, recoil and pointing limits, disposal requirements and provider safety process. The parametric recoil cases in Table IV must not be read as compatibility findings.

C. Integration Interfaces

A candidate integration would need a provider-approved mechanical interface, an electrical recharge and inhibit interface, an authenticated command/telemetry interface, and verified attitude, thermal, magnetic, debris and disposal constraints. The current CAD includes only a reference ring shape and bracket assumption; it does not establish ESPA or ESPA-Grande accommodation. An assumed 150–300 W recharge allocation, state-vector input and shot authorization are model inputs rather than approved host services. The deployer would command the modeled release only inside a host-authorized window, and the host would retain attitude and disposal authority. Payload-side exposure also matters: the modeled permanent-magnet field at the near payload face exceeds a representative magnetometer full scale, so compatibility requires a named payload and an explicit exclusion, standoff or shielding decision (Sec. XII).

A multi-shell mission could combine host manoeuvres and deployer release impulses, but its value must be calculated against spring release, differential drag, onboard propulsion and orbital transport with the same manifest and deadlines. A prior claim that release timing alone created persistent phase has been withdrawn. No complete twelve-payload campaign closes in the six sampled finite-burn cases, so this paper does not claim a delivered four-shell or twelve-target service.

D. Position in the Deployment Landscape

Spring deployers provide a mature reference for low-velocity separation; a host manoeuvre, differential drag or an orbital transfer service can add mission-level placement. VOLLEY is an unbuilt computational candidate for commanded release. Table V contrasts the studied functions and flags where

TABLE V

SERVICE TIERS FOR SMALL-SATELLITE ORBIT PLACEMENT. VOLLEY'S SYSTEM MASS IS 126.6 KG DRY AND 174.6 KG LOADED WITH TWELVE 3U SATELLITES

Attribute	Spring deployer	VOLLEY (this work)	Transfer vehicle
Release / transfer	low-m/s device-dependent	16.029 m/s model point	mission-dependent
State targeting	impulse and host-dependent	0.0274 m/s simulated 3σ	service-dependent
Payload interface	qualified per device	compatibility unverified	service-dependent
System mass, 12×3U	~72 kg comparator	126.6 kg dry model	quotation needed
Plane/LTAN change	host/drag-dependent	not demonstrated	service-dependent
Evidence status	flight heritage	computational study	provider-specific

a common-mission comparison or supplier-backed mass is absent. The approximate 6 kg spring canister used elsewhere in this paper implies about 72 kg for twelve 3U units, while the modeled Gen5 dry mass is 126.6 kg. No gap in the commercial market is established by this technical comparison.

E. Demonstration Path

F. What Air Costs a Ground Test

The machine operates in vacuum and every velocity above is computed with no aerodynamic term, but the full-scale demonstration fires a mass simulator down the track in a laboratory at sea level, so the comparison that test exists to make is between a measured velocity and a computed one that omits the air. A steady RANS solution ($k-\omega$ SST) around the sled and payload *as generated*, meshed directly from the CAD surfaces rather than from an idealised body, gives a drag coefficient of 0.52 on a 201.6 cm² projected frontal area at $Re = 7.0 \times 10^5$, i.e. 1.73 N at exit velocity. Because the profile is position-scheduled, v^2 rises linearly with distance and so does the drag, making the work over the 1.30 m acceleration zone 1.13 J and the exit-velocity deficit 5.1 mm/s, or 0.031 % (Fig. 19).

That figure is small against the design point and *not* small against the quantity a ground test resolves: it is 19 % of the 3σ closed-loop dispersion. An atmospheric ground test can be considered for initial velocity measurements, but its air-load correction and uncertainty must be reported; vacuum and thermal testing remain separate qualification questions. The distinction depends on how the test is run: open-loop, the deficit appears as 5.1 mm/s of velocity; closed-loop, the servo nulls it and it appears instead as 1.73 N of additional commanded force, 0.125 % of the shot.

Two properties of the solution are reported rather than presented as convergence. The velocity residuals fall three orders and then plateau near 5×10^{-2} , which is what a steady solver does on a massively separated wake, so the force is a mean over a window of iterations quoted with its spread, 1.734 ± 0.144 N. And the viscous contribution is a turbulent flat-plate bound over the meshed wetted area rather than a solved wall shear stress, 17.7 % of the total; the pressure term, which dominates, is integrated directly from the solved field over the body surface. Adding the stator plates *reduces* drag by 13 % and halves the oscillation, the plate acting as a wake splitter rather than as a confining wall, so the free-stream figure quoted above is the conservative one.




Fig. 17. Mid-plane slice ($z = 0$) of the converged fine-mesh solution: kinematic pressure above, speed below, flow left to right. Stagnation on the forward face, separation at the shoulders, and a wake that has not recovered by $x = 2.2$ m. The pressure term shown here is integrated directly from the solved field; the viscous term is not solved and is bounded by a turbulent flat-plate correlation over the meshed wetted area, as the text states.

A future instrumented ground article could measure sled and payload velocity, feed forces, tip-off, brake temperature and cycle reliability against predeclared bands. Its length, containment, cost and site approvals require engineering and a quotation; this study supplies neither. A captive hosted test would require a named platform, provider interface and safety approval even if the payload were never released.

VII. SENSITIVITY AND UNCERTAINTY

Because several inputs carry manufacturing or environmental spread, the rated point is bracketed by a first-order sensitivity analysis (Table VI). Each coefficient is the local partial derivative of exit velocity or thrust with respect to the listed variable, evaluated at the rated point. The dominant contributors are payload mass and magnet remanence; the velocity model is mild in all of them, and the closed-loop servo further collapses the delivered dispersion to 0.0274 m/s (Sec. V-A) because it commands current against measured position rather than trusting the plant. Root-sum-square combination of the manufacturing terms (magnet grade $\pm 2\%$, gap ± 0.05 mm, resistance $\pm 5\%$, ESR $\pm 20\%$) gives an open-loop velocity spread of $\pm 2.1\%$, consistent with the Monte Carlo of Fig. 9. Environmental terms are handled by rating margin: the N45SH remanence temperature coefficient of $-0.11\%/K$ over a ± 40 K operating band moves the thrust constant by $\mp 4.4\%$, so the worst-case cold-to-hot requirement of 132 kA/m stays inside the 140 kA/m current rating without re-tuning. Atmospheric and solar-activity spread do not affect the machine and enter only through orbital lifetime, where Sec. V-B records that the sweep previously offered as evidence of invariance could not, by construction, have detected a change.




Fig. 18. Force history and mesh convergence for the same case. The force is a mean over a window of iterations quoted with its spread, 1.734 ± 0.144 N, because a steady solver on a massively separated wake plateaus rather than converges; $6.3\times$ the cells moves the answer by 4.86 %.




Fig. 19. Air drag along the acceleration zone, and the resulting exit-velocity deficit set against the design point and against the dispersion a ground test exists to resolve.

TABLE VI
FIRST-ORDER SENSITIVITIES AT THE RATED POINT

Variable	Affects	Sensitivity
Payload mass	v	$-\frac{1}{2}$ (%/%) of m_{tot}
Magnet remanence B_r	F, v	$+1.0, +0.5$ %/%
Current density	F, v	$+1.0, +0.5$ %/%
Winding resistance	η	-0.15 %/%
Capacitor ESR	sag	$+0.04$ V/mΩ
Gap asymmetry	K_t	-13 %/mm
Sensor noise (8 mm/s)	dispersion	sets 3σ floor
Temperature (B_r)	K_t	-0.11 %/K

VIII. MECHANICAL ANALYSIS

The structural load path is sized to the two governing cases: launch vibration during ascent and the arrest impulse during operation. Fastener and pin margins are quoted on ultimate strength at a 1.4 design factor per common launch-hardware practice, consistent with the qualification philosophy of GEVS [21] and ECSS structural standards [22].

Track stiffness. The dominant flight requirement is a first natural frequency above the launch-vehicle primary band, conventionally 70–100 Hz for secondary structure. Modeling the twin-longeron track as a uniform beam carrying the 20 kg stator distributed mass, the fundamental is $f_1 =$

$(\lambda_1^2/2\pi L^2)\sqrt{EI/\mu}$ with EI from two 7075 box sections. A simply-supported idealization gives 48 Hz, but bracketing both ends into the ESPA ring (fixed-fixed) raises the first mode to 109 Hz, clearing the requirement; the design therefore specifies end-fixed mounting with an intermediate rib.

Inter-array attraction. The two Halbach faces attract across the gap with a Maxwell stress $p = B^2/2\mu_0 \approx 120$ kPa at the 0.55 T mean working field, or 3.7 kN over the array footprint. Carried as a bending load into the titanium side plates, this produces 33 MPa against the 880 MPa Ti-6Al-4V yield, a margin above 20; the attraction is structurally minor and, being static, does not fatigue-cycle. The flat-plate expression overstates the force: a three-dimensional integration of the field gradient over the meshed blocks returns 2.69 kN, some 37% lower, because the assumed 0.55 T is not the mean normal field over the separation plane. Jensen's inequality acts in the opposite direction and does not explain the overestimate. Extending the Maxwell-stress plane beyond the magnet footprint gives 2.680 kN, within 0.25% of the 2.687 kN volume-force calculation; this is model-to-model agreement, not experimental validation. The earlier finite-element load remains conservative.

Arrest and roller loads. At the 200 g arrest cap the 9.445 kg sled develops 18.5 kN axially, reacted by the ring-spring stop; the pitch couple during braking loads the guide rollers to roughly 1.48 kN per pair. This is the load path most affected by adopting the CAD-derived sled mass, it very nearly doubles, and it is the one place where the heavier sled makes the machine harder rather than merely slower. Bearing selection at 1.48 kN per pair is within the rating of a 440C/Si₃N₄ hybrid of this size, but the reuse duty is not addressed here: the sled is braked twelve times per campaign in vacuum, and no lubrication or contact-life assessment has been performed. The abort latch, engaged only for a commanded abort, carries the 13.4 kg carriage at the abort deceleration, 0.42 kN, through dual A-286 pins at a margin above eight.

Fatigue and tolerance. Operational cycling is benign: the magnet bond sees 0.12 MPa per shot against a 10 MPa

allowable, far below any endurance concern over the few hundred shots of a service life. Positional accuracy is set by the winding-to-magnet gap, whose $-13\%/mm$ thrust sensitivity dictates a ± 0.05 mm shim tolerance to hold thrust within $\pm 0.65\%$; this is a standard precision-shim assembly, not a novel requirement.

A. Architecture reliability, and the bar it sets

A spring dispenser is twelve independent one-shot mechanisms in parallel: one failure costs one satellite. This design is one mechanism in series with itself, cycled twelve times. The sled, stator, converter, energy store, sequencer and brake serve every shot, and the escapement and retention gate cycle twelve times each, so a failure at shot k forfeits shots k through twelve. Enumerating the shot chain, nine of thirteen elements forfeit the remaining manifest on a single failure, against zero for a spring dispenser, and nine shared elements over twelve cycles is 108 opportunities to fail. This is structural and no component quality removes it.

What component quality does change is where the comparison lands. Writing p for the probability that the shared chain survives one shot, expected delivery is $\sum_{k=1}^{12} p^k$, which departs from a spring's independent $12q$ quickly. Against a spring at $q = 0.99$, matching on *satellite count* requires $p = 0.9985$, which is not a realistic target for a twelve-cycle electromechanical system without flight heritage. Matching on *delivered orbital life* requires only $p = 0.9347$, because each delivered satellite is worth 1.50 times a spring-deployed one. Expressed per element per cycle, the requirement is $r \geq 0.99326$.

That is the first quantitative reliability requirement this design has carried, and it is not yet met or measured: no cycle-life test exists for the escapement, the gate or the sled. Two candidate mitigations are defined but are deliberately not credited above, because crediting them without analysis is the optimism the exercise exists to avoid. Retention gates are per cassette, so one gate fault forfeits six rather than twelve. And the winding is segmented: a dead segment is a length the sled coasts over rather than a stopped machine, so at four segments a failure downstream of the breech still exits at 14.19 m/s, 86.6% of nominal. The exception is the breech segment itself, where no force acts on a stationary sled and the shot never starts, which makes duplicating or overlapping that one segment the cheapest available reliability action.

IX. THERMAL ANALYSIS

The shot model assigns approximately 855 J to stator copper, 1162 J to sled braking, 93 J to the converter and 78 J to the bank ESR per rated cycle. Multiplying an assumed twelve-shot total of 26.6 kJ by a uniform 13.5 kJ/K structure thermal capacitance gives a 2.0 K *bulk-average* rise; it says little about local peak temperatures. The corrected 0.344 kg fin mass gives a 9.0 K adiabatic rise per shot and a 108 K twelve-shot no-cooling bound. A transient conduction/radiation model has not established cooling between shots at the adopted 1200 s cadence. The 0.32 m² radiator estimate uses an assumed

130 W rejection case and is not a closed thermal design. Component survival temperatures, heater power and irreversible magnet loss require supplier data and orbit-specific hot/cold cases. The rated thermal claims remain model estimates, not environmental qualification.

X. POWER ELECTRONICS

The electrical architecture proposes a 96 V, thirty-two-cell pulse bank and a three-phase SiC inverter with position-synchronized control. The winding model predicts 19.7 μ H, a 373 A peak phase current and 25.2 m Ω per phase at the rated point; 320 A is the reported DC-link peak and is a different quantity. At the assumed 20–40 kHz switching range, the calculated current ripple and copper contribution are model outputs. Device ratings, switching overshoot, gate-drive loss, cell balancing, fault-clearing energy and flight derating have not been checked with a selected part list or a physical converter.

Current limiting, bus clamping, desaturation shutdown, balancing, bleed/passivation, segmentation, shielding and dual sensing are proposed controls. None has a closed circuit schematic, thermal qualification, EMC measurement, fault-injection result or provider acceptance in this study. A shorted segment can halt a shot or create an unsafe transient; segmentation cannot yet be credited as fail-operational. The permanent-magnet field and inverter transients may reach other payloads, so scheduling after primary separation is not a substitute for payload-specific magnetic and EMC analysis.

XI. RELIABILITY AND FAILURE MODES

Table VII summarizes the principal failure modes in an FMEA form, with the mitigation designed into the architecture. The intended safety goals are to retain the satellite through ascent, prevent unintended release after any credible fault and recover a safe state when possible. Single-fault tolerance and graceful degradation have not been established for the feeder, bank, inverter or brake.

XII. SPACE ENVIRONMENT, EMC, AND STANDARDS

LEO accommodation would require a selected material and component list checked for outgassing, atomic oxygen, thermal cycling, radiation, vibration and shock against the actual provider and mission standards [22]. The current model proposes low-outgassing polymers, protected surfaces and N45SH magnets, but no supplier qualification record, launch-load test or environmental test is supplied. The calculated 109 Hz first mode and selected load margins apply to stated structural idealizations; they do not demonstrate assembly survivability.

The permanent-magnet sled is a payload compatibility concern. The modelled back-face field is 22.7 mT at 10 mm, 4.3 mT at 20 mm and 0.4 mT at 50 mm for the stated geometry. Across a 3U payload envelope, the calculated field is 61.1 mT at the near face and 0.341 mT at the far face. The near value exceeds a representative ± 100 μ T magnetometer full scale by 611 times; even the far value exceeds it by

TABLE VII
PRINCIPAL FAILURE MODES AND MITIGATIONS

Failure	Effect	Mitigation
Sled jam	Manifest stranded	Caged escapement; controlled-disposal host. The dual-cassette split bounds an <i>escapement</i> or <i>gate</i> fault to half the manifest, but not a sled fault: one sled serves all twelve cells
Latch fails	Early release	Latch loaded only on abort; inertia retains during accel; three-inhibit fire chain
Sensor fault	Wrong velocity	Redundant seat switches, encoder, photogate cross-check; no-fire on disagreement
Capacitor cell	Reduced energy	32 series cells; one shorted cell drops rail 3%, within margin
Inverter device	Phase loss	Segmented winding; dual gate drive; degraded-thrust continue
Magnet debond	Field loss	84× bond margin at 200 g; taper-limited arrest
CubeSat misalign	Feed jam / tip-off	Seat switches + optical check; no-fire interlock
Brake fault	Sled overrun	Ring-spring hard stop backs eddy brake; motor regen trim
Controller fault	Loss of shot	Safe-state inhibits; watchdog; abort to 45% stroke

3.4 times. These are model results, not a qualified payload field map. No published VOLLEY ICD defines a keep-out radius, and the proposed steel septa and shielding have not been evaluated as installed hardware. Magnetometer recovery, magnetorquer operation, permanent magnetisation of payload parts and switching-transient coupling must be assessed with a named satellite. Timing a firing after primary separation does not remove the static-field exposure of stored spacecraft.

XIII. MANUFACTURING, COST, AND SCALABILITY

A possible fabrication sequence uses a wound copper stator on an insulating former, a restrained fixture for assembly of the magnet arrays, and late loading of spacecraft after mechanism inspection. Proposed slot and gap tolerances of ± 0.1 and ± 0.05 mm are design targets, not supplier capabilities or a completed stack-up. The current STEP shapes omit fasteners, harness, joints and feed motion, and the side-fed reference assembly clashes; therefore no released manufacturing drawing set, assembly process or recurring maintenance interval can be claimed. An eventual ground article would require measured current-to-velocity calibration and acceptance of the complete release cycle.

The cost model is parametric and has no supplier quotations. In its assumed bill of materials, avionics, pulse storage and sensing dominate the estimated recurring hardware cost more than the magnet set; this ordering is conditional on unquoted prices and the incomplete installed system. The rated shot draws 2.78 kJ, less than 1 Wh of electrical energy, but that is not a service-cost comparison: host power, storage, integration, qualification, mass, reliability, operations and any host propellant must be counted. A matched mission comparison with qualified spring systems, host manoeuvres, onboard propulsion and transport services remains open.

Ideal scaling trends ($v \propto \sqrt{L}$ at fixed acceleration) motivate a wider trade, but the payload-family table is model output, not accommodation of other CubeSat sizes. More payloads change feeder geometry, structure, reliability and host burden; the current twelve-slot reference already fails a width check. Lunar or deep-space use would require new interfaces, thermal/radiation design, mission dynamics and payload qualification, and is outside this study.

XIV. OPERATIONAL CONCEPT AND MATURATION

The proposed operations sequence is launch with the deployer dormant; host and payload health checks; bank charge inside a provider-authorized window; one commanded release followed by independent confirmation of separation and safe sled arrest; and repetition only after fault review. Inhibit logic, an early-stroke abort window, bank passivation and eventual host disposal are design requirements, not demonstrated functions. POEM-3 provides a hosted-stage precedent [14], but no POEM or other provider has approved this sequence for VOLLEY.

The next evidence sequence is a reviewed feeder re-design, component-level bench tests, an instrumented integrated ground article, environmental tests against a named payload and host, and only then a provider-approved captive or deployment demonstration. Each stage needs predeclared acceptance bands and a decision review. No technology-readiness level is assigned from the present numerical results, and a future test does not automatically confer one.

XV. LIMITATIONS

The new 20-instance FreeCAD reference assembly is reviewable as native FCStd and STEP, but its as-drawn side-fed placement fails the 11 mm width screen and has exact track-cassette overlap (Fig. 20). It is not a solved feeder, tolerance stack or manufacturing release. Eight source solids are CadQuery-generated and re-exported through FreeCAD; their native FreeCAD instances do not carry parametric feature histories. Seven limits bound the claims. The first is that payload structural compatibility is not established, and cannot be established by this work alone. The 25 g ceiling this machine is designed to is a requirement it sets on itself. It was previously presented as following from the CubeSat qualification environment, by integrating a GEVS random-vibration spectrum to 14.1 g_{rms} and taking three times that as a bounding load; that reasoning is withdrawn, because a root-mean-square level over a broadband random spectrum is not a quasi-static limit load and a 3σ multiple of it is a peak-response estimate for that spectrum rather than a structural capability. No universal quasi-static qualification level for CubeSats exists to replace it: the CubeSat Design Specification [23] defers test levels to the launch provider, so they vary by vehicle and mission. The defensible statement is therefore narrow: the rated design acceleration is 10.1 g for 162 ms, and whether any specific satellite tolerates it depends on that satellite's own qualified load environment and on an integration review that has not been performed. No payload




Fig. 20. Gen5 side-fed reference-section failure. The two cassette widths and track require 537 mm within a 526 mm internal width before clearance. FreeCAD and CadQuery exact-solid intersections each find 32,915 mm³ overlap per side. A new feed or enclosure configuration would require renewed analysis.

has been mechanically qualified against this machine, and the payload-family accelerations in Table III, which run to 14.0 g at PocketQube, carry the same caveat. The atmosphere model is static, so absolute lifetimes carry severalfold uncertainty; an earlier version of this paper offered the invariance of the lifetime *ratio* as the defensible result in their place, and Sec. V-B records why that was wrong, the sweep supporting it could not have detected a change. The ratio is now quoted at a stated activity level and carries the same uncertainty as everything else derived from a static atmosphere. The thrust constant is winding-resolved in two dimensions; three-dimensional end effects of a few percent remain for finite-element closure. The brake law is first-order plate drag and remains an unverified arrest approximation. Conjunction figures, although refined at candidate minima, do not replace per-shot collision-avoidance products. Host budgets are parametric wherever stage properties are not public. The 9.445 kg sled input comes from historical Gen3 CAD solid volumes, whereas other dry-mass lines are modeled or assumed. This is not a mass extraction from a complete Gen5 assembly. The former 4.86 kg parametric sled estimate disagrees with those volumes. The CAD-derived value is adopted here and sets the headline velocity, but it is the as-drawn, unpocketed geometry, the structural analysis reports a 17-fold stress margin on the chassis plate, so a rib-stiffened redesign would recover mass, and none has been evaluated. The sensitivity, mechanical, thermal, power-electronics, and reliability analyses added here

(Secs. VII–XIV) address the principal engineering questions only at analysis level. Several remain feasibility-significant. The commercial capacitor strings examined do not meet the ESR requirement; a flywheel motor-generator has since been shown to clear it at 35 m Ω against the 68 m Ω ceiling, needing no architecture change, but at mass parity rather than saving. The three-dimensional depth correction has been applied to the analytic field model and the result was unfavourable: resolving the field across the array’s 90 mm depth, rather than sampling it at the centre plane and multiplying, reduces K_t by 4.42% to 10.54 N per kA/m and v_{exit} to 16.03 m/s. That correction was computed on 2026-08-10 and held for three days rather than applied, because moving a frozen baseline is not a verification run’s decision to take; it was adopted under change control on 2026-08-13, on the grounds that depth-resolving the integral is a change to the physics rather than a pasted correction factor and that collapsing the depth quadrature to a single plane reproduces the superseded value exactly, which keeps the ratio auditable. Every figure in this paper is quoted after it. The three-dimensional solve run the same day supports the underlying field model, a reduced-scalar-potential magnetostatic PDE on a 315,370-node tetrahedral mesh agrees with it to 0.059% at midgap, but it tests the field, not the depth average, and no three-dimensional check of the integrated thrust constant itself exists. The earlier 11.03 N per kA/m and 16.39 m/s appear here only where a superseded value is named as such. And hardware validation has not begun.

XVI. CONCLUSION

Gen5 is a fixed, system-level computational design for a magazine-fed linear-motor deployer. In the stated model it ejects an unmodified 3U CubeSat at 16.0 m/s and 10.1 g , with 18.8% net electrical-to-payload efficiency and 0.0274 m/s simulated dispersion; the dry mass rollup is 126.6 kg. The orbital model predicts a 28.8 km semi-major-axis increase and a 1.60 lifetime multiplier at mean solar activity for one prograde shot. The commanded release profile is the proposed capability. The model also fails its mass-per-satellite criterion, and its side-fed reference CAD assembly has a measured track–cassette clash. It has no measured launch, payload qualification, selected host or completed integration review. The result is an analysed design with explicit limits, not a flight-qualified product. A matched mission comparison and physical tests are needed before claiming deployment or economic advantage.

XVII. FUTURE WORK: GEN6 SCALE-UP

Gen5 remains the sole design and performance baseline reported here. Gen6 is the proposed follow-on research programme: extend individually commanded orbital release toward a 1 km/s-class exit velocity and, if justified by mission analysis, beyond. This is a target rather than a selected configuration or a demonstrated result. At constant acceleration, 1 km/s requires about 2.0 km of stroke at 25 g or 0.51 km at 100 g from rest. Achieving it within Gen5’s 1.3 m acceleration stroke would require approximately $3.9 \times 10^4 g$ on an ideal constant-acceleration estimate. The payload class and qualified

load limit, guide geometry, energy and heat rejection, recoil and host control authority, containment, release interface and mission value must therefore be re-established together. Earlier gas-guide and independent spring-cell studies are unselected comparisons, not validated Gen6 hardware. The next evidence gates are a specified payload and host, a complete installed-system trade, an independently checked dynamics model and calibrated hardware tests. None of those gates is closed by the Gen5 calculations.

REFERENCES

- [1] C. Lewis, "Failure of the ball-lock mechanism on the NanoRacks CubeSat deployer," in *Proc. 44th Aerospace Mechanisms Symp.*, 2018.
- [2] JAXA, "JEM payload accommodation handbook, vol. 8: Small satellite deployment interface control document," via UNOOSA KiboCUBE.
- [3] NASA SSVI, "State-of-the-art of small spacecraft technology: Integration, launch, deployment, and orbital transport."
- [4] B. N. Turman *et al.*, "Coilgun launcher for nanosatellites," Sandia National Laboratories, OSTI.
- [5] R. J. Kaye *et al.*, "Electromagnetic coilgun launcher for space applications," Sandia National Laboratories, OSTI 125180.
- [6] K. Halbach, "Design of permanent multipole magnets with oriented rare earth cobalt material," *Nucl. Instrum. Methods*, vol. 169, pp. 1–10, 1980.
- [7] H. Q. Yamaguchi and M. Morishita, "End effects of Halbach field magnet-type linear motor," *Electrical Engineering in Japan*, 2023, doi:10.1002/eej.23450.
- [8] R. F. Post and D. D. Ryutov, "The Inductrack: A simpler approach to magnetic levitation," *IEEE Trans. Appl. Supercond.*, vol. 10, no. 1, pp. 901–904, 2000.
- [9] F. Faure *et al.*, "Qualification of COTS supercapacitors for space applications," *ESA Space Passive Component Days*, 2018.
- [10] "Towards supercapacitors in space applications," *European Space Power Conf.*, 2017.
- [11] E. Díez-Jiménez, C. Alén-Cordero, R. Alcover-Sánchez, and E. Corral-Abad, "Modelling and test of an integrated magnetic spring-eddy current damper for space applications," *Actuators*, vol. 10, no. 1, art. 8, 2021; together with commercial eddy-current damper flight-heritage documentation, not independently retrieved.
- [12] C. Foster *et al.*, "Constellation phasing with differential drag on Planet Labs satellites," *J. Spacecraft and Rockets*, vol. 55, no. 2, 2018.
- [13] C. Foster, H. Hallam, and J. Mason, "Orbit determination and differential-drag control of Planet Labs CubeSat constellations," arXiv:1509.03270, 2015.
- [14] ISRO, "PSLV-C58/XPoSat mission: POEM-3 accomplishes zero orbital debris mission," Mar. 2024.
- [15] ISRO, "POEM-4 completes 1000 orbits in space," 2025.
- [16] NanoRacks, "NanoRacks CubeSat deployer interface definition document," 2018.
- [17] D. A. Vallado, *Fundamentals of Astrodynamics and Applications*, 4th ed. Microcosm Press, 2013.
- [18] Skyroot Aerospace, "Orbit Adjustment Module full stage-level test," public statement, Oct. 14, 2025.
- [19] J. Foust, "Skyroot prepares for first orbital launch attempt," *SpaceNews*, Jul. 7, 2026.
- [20] Exolaunch, "CarboNIX" and "EXOpod Nova" product documentation; Exolaunch–Skyroot partnership announcement, Oct. 14, 2025.
- [21] NASA, "General Environmental Verification Standard (GEVS)," GSFC-STD-7000A, 2013.
- [22] ECSS, "Space engineering: Structural and mechanical / materials standards," ECSS-E-ST-32 and ECSS-Q-ST-70 series.
- [23] California Polytechnic State Univ., "CubeSat Design Specification," Rev. 14, 2020.
- [24] E. Kulu, "Nanosats Database," nanosats.eu, counts as of Jan. 1, 2026.
- [25] Novaspaces, "Prospects for the Small Satellite Market," 11th ed., 2026. Market forecast, not a measurement.
- [26] Bellatrix Aerospace, "Bellatrix Aerospace and NSIL sign MoU to integrate Bellatrix's OTV in NSIL's launch missions," Oct. 9, 2024.
- [27] H. Feng, Y. Yang, and Y. Wu, "Design and reachable domain analysis of on-orbit electromagnetic launcher for CubeSats," *Int. J. Aerospace Eng.*, vol. 2025, art. 3000765, 2025.
- [28] D. Xu, H. Yue, Y. Zhao, F. Yang, J. Wu, X. Pan, T. Tang, and Y. Zhang, "Improved A* algorithm for path planning based on CubeSats in-orbit electromagnetic transfer system," *Aerospace*, vol. 11, no. 5, art. 394, 2024.
- [29] Y. Zhao, Y. Zhang, Z. Zhao, C. Li, L. Zhang, X. Yang, H. Yue, *et al.*, "Simulation analysis and experimental verification of the transport characteristics of a high-volume CubeSat storage device," *Aerospace*, vol. 12, no. 6, art. 466, 2025.
- [30] Y. Zhao, H. Yue, X. Mu, X. Yang, and F. Yang, "Design and analysis of a new deployer for the in-orbit release of multiple stacked CubeSats," *Remote Sensing*, vol. 14, no. 17, art. 4205, 2022.
- [31] M. Einat and Y. Orbach, "A multi-stage 130 m/s reluctance linear electromagnetic launcher," *Sci. Rep.*, vol. 13, 2023.